

Regression Model

Supervised learning has two type of model:

- 1)Regression
- 2)classification

1. Regression (CH-5)

Regression is continuous:

- > LR(linear regression)
- > Polynomial Regression

2. Classification (CH-6)

Classification is discrete:

- > KNN
- > Decision Tree

LR(linear regression)

simple linear regression(SLR):

$$y=mx+c$$

where, y=dependent variable

x=independent variable

m=co-efficient

c=intercept

multiple linear regression(MLR):

Step of Supervised learning process-

step1 : Define the problem

step2 : Collect and Prepare data

step3 : Data preprocessing

step4 : Feature Engineering

step5 : Select a model

step6 : Split the data in test/train

step7 : Train the model

step8 : Validate hyper parameters

step9 : Validate the model on the basis of tesing data

step10: Iterate the model for improvement

step11: Deploy a model

Simple linear regression with the help of excel

Simple linear regression model

```
In [5]: import pandas as pd
import numpy as np

df=pd.read_csv("placement_03-06.csv")
df
```

Out[5]:

	cgpa	package
0	6.89	3.26
1	5.12	1.98
2	7.82	3.25
3	7.42	3.67
4	6.94	3.57
...
195	6.93	2.46
196	5.89	2.57
197	7.21	3.24
198	7.63	3.96
199	6.22	2.33

200 rows × 2 columns

```
In [6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   cgpa        200 non-null    float64
1   package     200 non-null    float64
dtypes: float64(2)
memory usage: 3.2 KB
```

```
In [7]: from sklearn.model_selection import train_test_split  
y=df['package']  
y
```

```
Out[7]: 0      3.26  
1      1.98  
2      3.25  
3      3.67  
4      3.57  
...  
195    2.46  
196    2.57  
197    3.24  
198    3.96  
199    2.33  
Name: package, Length: 200, dtype: float64
```

```
In [11]: x=df.drop('package',axis=1)  
x
```

```
Out[11]:
```

	cgpa
0	6.89
1	5.12
2	7.82
3	7.42
4	6.94
...	...
195	6.93
196	5.89
197	7.21
198	7.63
199	6.22

200 rows × 1 columns

```
In [12]: type(y)
```

```
Out[12]: pandas.core.series.Series
```

```
In [13]: type(x)
```

```
Out[13]: pandas.core.frame.DataFrame
```

```
In [14]: y.shape
```

```
Out[14]: (200,)
```

```
In [15]: x.shape
```

```
Out[15]: (200, 1)
```

```
In [49]: x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.8,random_state=
```

```
In [50]: print(y_train.shape)
print(y_test.shape)
print(x_train.shape)
print(x_test.shape)
```

```
(160,)
(40,)
(160, 1)
(40, 1)
```

```
In [51]: from sklearn.linear_model import LinearRegression
lr=LinearRegression()
model=lr.fit(x_train,y_train)
```

```
In [52]: print("m = ",model.coef_)
```

```
m = [0.58154877]
```

```
In [53]: print("c =" ,model.intercept_)
```

```
c = -1.0859839580358024
```

```
In [54]: # y=mx+c
y_pred=model.predict(x_test)
print(y_pred)
```

```
[2.9383335  4.36894346 3.18258398 1.89736121 3.49662031 3.35123312
 2.76968435 2.94996447 3.07208971 3.94441286 3.57222165 2.94996447
 2.75805338 2.64755911 3.67108494 3.2174769  3.97930579 2.90925606
 2.19395108 3.31052471 4.29915761 2.8918096  1.87409926 2.30444534
 3.62456104 2.12998071 3.9269664  2.36841571 1.5716939  2.06601035
 2.31026083 3.6885314  3.5024358  3.03719679 2.57195777 2.39167766
 3.170953   3.82228762 3.15932203 2.94414898]
```

```
In [55]: # y_pred-y_test  
diff=pd.DataFrame({"actual":y_test,  
                  "predicted":y_pred})  
diff
```

Out[55]:

	actual	predicted
58	3.09	2.938333
40	4.02	4.368943
34	3.42	3.182584
102	1.37	1.897361
184	3.14	3.496620
198	3.96	3.351233
95	2.79	2.769684
4	3.57	2.949964
29	3.49	3.072090
168	3.52	3.944413
171	3.76	3.572222
18	2.98	2.949964
11	2.60	2.758053
89	2.72	2.647559
110	3.76	3.671085
118	2.88	3.217477
159	4.08	3.979306
35	2.87	2.909256
136	2.10	2.193951
59	3.31	3.310525
51	3.79	4.299158
16	2.35	2.891810
44	1.86	1.874099
94	2.42	2.304445
31	3.89	3.624561
162	2.55	2.129981
38	4.36	3.926966
28	2.24	2.368416
193	1.94	1.571694
27	2.16	2.066010
47	3.26	2.310261
165	4.08	3.688531
194	3.67	3.502436
177	3.64	3.037197
176	3.23	2.571958

	actual	predicted
97	2.84	2.391678
174	2.99	3.170953
73	4.03	3.822288
69	2.94	3.159322
172	2.51	2.944149

```
In [56]: # (mean absolute error)
# mae=1/n |y_act-y_pred|
```

```
In [57]: from sklearn.metrics import mean_absolute_error
mae=mean_absolute_error(y_test,y_pred)
print(mae)
```

0.29931188593316804

```
In [58]: from sklearn.metrics import mean_squared_error
mse=mean_squared_error(y_test,y_pred)
print(mse)
```

0.1370062519255722

```
In [59]: from sklearn.metrics import r2_score
rs=r2_score(y_test,y_pred)
print(rs)
```

0.7283345498058083

Multiple linear regression model-1

```
In [67]: import pandas as pd
df=pd.read_csv("Advertising_03-06.csv")
df
```

Out[67]:

	Unnamed: 0	TV	Radio	Newspaper	Sales
0	1	230.1	37.8	69.2	22.1
1	2	44.5	39.3	45.1	10.4
2	3	17.2	45.9	69.3	9.3
3	4	151.5	41.3	58.5	18.5
4	5	180.8	10.8	58.4	12.9
...
195	196	38.2	3.7	13.8	7.6
196	197	94.2	4.9	8.1	9.7
197	198	177.0	9.3	6.4	12.8
198	199	283.6	42.0	66.2	25.5
199	200	232.1	8.6	8.7	13.4

200 rows × 5 columns

```
In [68]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 5 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Unnamed: 0  200 non-null    int64
1   TV          200 non-null    float64
2   Radio       200 non-null    float64
3   Newspaper   200 non-null    float64
4   Sales       200 non-null    float64
dtypes: float64(4), int64(1)
memory usage: 7.9 KB
```



```
In [69]: df.drop("Unnamed: 0",axis=1,inplace=True)
df
```

Out[69]:

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	9.3
3	151.5	41.3	58.5	18.5
4	180.8	10.8	58.4	12.9
...
195	38.2	3.7	13.8	7.6
196	94.2	4.9	8.1	9.7
197	177.0	9.3	6.4	12.8
198	283.6	42.0	66.2	25.5
199	232.1	8.6	8.7	13.4

200 rows × 4 columns

```
In [70]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0    TV          200 non-null    float64
1    Radio       200 non-null    float64
2    Newspaper   200 non-null    float64
3    Sales       200 non-null    float64
dtypes: float64(4)
memory usage: 6.4 KB
```

```
In [71]: from sklearn.model_selection import train_test_split
y=df['Sales']
y
```

Out[71]:

0	22.1
1	10.4
2	9.3
3	18.5
4	12.9
...	
195	7.6
196	9.7
197	12.8
198	25.5
199	13.4

Name: Sales, Length: 200, dtype: float64

```
In [72]: x=df.drop('Sales',axis=1)
x
```

Out[72]:

	TV	Radio	Newspaper
0	230.1	37.8	69.2
1	44.5	39.3	45.1
2	17.2	45.9	69.3
3	151.5	41.3	58.5
4	180.8	10.8	58.4
...
195	38.2	3.7	13.8
196	94.2	4.9	8.1
197	177.0	9.3	6.4
198	283.6	42.0	66.2
199	232.1	8.6	8.7

200 rows × 3 columns

```
In [73]: print(type(x))
print(type(y))

<class 'pandas.core.frame.DataFrame'>
<class 'pandas.core.series.Series'>
```

```
In [74]: print(x.shape)
print(y.shape)

(200, 3)
(200,)
```

```
In [77]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.15,random_state=
```

```
In [78]: print(x_train.shape)
print(x_test.shape)
print(y_train.shape)
print(y_test.shape)

(170, 3)
(30, 3)
(170,)
(30,)
```

```
In [80]: from sklearn.linear_model import LinearRegression
lr=LinearRegression()
model=lr.fit(x_train,y_train)
print("m =" ,model.coef_)
print("c =" ,model.intercept_)
```

```
m = [ 0.04666462  0.18481678 -0.00121229]
c = 2.898009326357265
```

```
In [81]: y_pred=model.predict(x_test)
print(y_pred)
```

```
[21.8561195  16.43070013  7.61358579 17.81414167 18.64171793 23.81414579
 16.29442126 13.26019548  9.10013819 17.24141586 14.3795469   9.89375986
 17.34765633 16.79371444 14.88188303 15.48747719 12.40242285 17.2108581
 11.28920355 18.17106497  9.35301379 12.68796292  8.76659008 10.48401019
 11.33546207 15.00377232  9.8013108  19.48893945 18.43960804 17.16086278]
```

```
In [83]: diff=pd.DataFrame({"actual":y_test,"predicted":y_pred})  
diff
```

Out[83]:

	actual	predicted
58	23.8	21.856120
40	16.6	16.430700
34	9.5	7.613586
102	14.8	17.814142
184	17.6	18.641718
198	25.5	23.814146
95	16.9	16.294421
4	12.9	13.260195
29	10.5	9.100138
168	17.1	17.241416
171	14.5	14.379547
18	11.3	9.893760
11	17.4	17.347656
89	16.7	16.793714
110	13.4	14.881883
118	15.9	15.487477
159	12.9	12.402423
35	12.8	17.210858
136	9.5	11.289204
59	18.4	18.171065
51	10.7	9.353014
16	12.5	12.687963
44	8.5	8.766590
94	11.5	10.484010
31	11.9	11.335462
162	14.9	15.003772
38	10.1	9.801311
28	18.9	19.488939
193	19.6	18.439608
27	15.9	17.160863

```
In [84]: from sklearn.metrics import mean_absolute_error
mae=mean_absolute_error(y_test,y_pred)
print(mae)
```

0.984560464503767

```
In [85]: from sklearn.metrics import mean_squared_error
mse=mean_squared_error(y_test,y_pred)
print(mse)
```

1.8945245763596135

```
In [86]: from sklearn.metrics import r2_score
rs=r2_score(y_test,y_pred)
print(rs)
```

0.882855181551423

Multiple linear regression model-2

```
In [91]: import pandas as pd
df=pd.read_csv("insurance_03-06.csv")
df
```

Out[91]:

	age	sex	bmi	children	smoker	region	expenses
0	19	female	27.9	0	yes	southwest	16884.92
1	18	male	33.8	1	no	southeast	1725.55
2	28	male	33.0	3	no	southeast	4449.46
3	33	male	22.7	0	no	northwest	21984.47
4	32	male	28.9	0	no	northwest	3866.86
...
1333	50	male	31.0	3	no	northwest	10600.55
1334	18	female	31.9	0	no	northeast	2205.98
1335	18	female	36.9	0	no	southeast	1629.83
1336	21	female	25.8	0	no	southwest	2007.95
1337	61	female	29.1	0	yes	northwest	29141.36

1338 rows × 7 columns

In [92]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   age         1338 non-null   int64
 1   sex         1338 non-null   object
 2   bmi         1338 non-null   float64
 3   children    1338 non-null   int64
 4   smoker      1338 non-null   object
 5   region      1338 non-null   object
 6   expenses    1338 non-null   float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
```

In [93]: `df.drop(["sex", "children", "region"], axis=1, inplace=True)`
`df`

Out[93]:

	age	bmi	smoker	expenses
0	19	27.9	yes	16884.92
1	18	33.8	no	1725.55
2	28	33.0	no	4449.46
3	33	22.7	no	21984.47
4	32	28.9	no	3866.86
...
1333	50	31.0	no	10600.55
1334	18	31.9	no	2205.98
1335	18	36.9	no	1629.83
1336	21	25.8	no	2007.95
1337	61	29.1	yes	29141.36

1338 rows × 4 columns

```
In [97]: df1=pd.get_dummies(df,drop_first=True)
df1
```

Out[97]:

	age	bmi	expenses	smoker_yes
0	19	27.9	16884.92	1
1	18	33.8	1725.55	0
2	28	33.0	4449.46	0
3	33	22.7	21984.47	0
4	32	28.9	3866.86	0
...
1333	50	31.0	10600.55	0
1334	18	31.9	2205.98	0
1335	18	36.9	1629.83	0
1336	21	25.8	2007.95	0
1337	61	29.1	29141.36	1

1338 rows × 4 columns

```
In [103]: from sklearn.model_selection import train_test_split
y=df1['expenses']
y
```

Out[103]:

0	16884.92
1	1725.55
2	4449.46
3	21984.47
4	3866.86
...	...
1333	10600.55
1334	2205.98
1335	1629.83
1336	2007.95
1337	29141.36

Name: expenses, Length: 1338, dtype: float64

```
In [104]: x=df1.drop('expenses',axis=1)
x
```

Out[104]:

	age	bmi	smoker_yes
0	19	27.9	1
1	18	33.8	0
2	28	33.0	0
3	33	22.7	0
4	32	28.9	0
...
1333	50	31.0	0
1334	18	31.9	0
1335	18	36.9	0
1336	21	25.8	0
1337	61	29.1	1

1338 rows × 3 columns

```
In [105]: print(type(x))
print(type(y))

<class 'pandas.core.frame.DataFrame'>
<class 'pandas.core.series.Series'>
```

```
In [106]: print(x.shape)
print(y.shape)

(1338, 3)
(1338,)
```

```
In [110]: x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.85,random_state=42)
print(x_train.shape)
print(x_test.shape)
print(y_train.shape)
print(y_test.shape)

(1137, 3)
(201, 3)
(1137,)
(201,)
```



```
In [111]: from sklearn.linear_model import LinearRegression
lr=LinearRegression()
model=lr.fit(x_train,y_train)
print("m =" ,model.coef_)
print("c =" ,model.intercept_)
```

```
m = [ 260.25098608  313.77333855 23758.11752343]
c = -11417.098577656794
```

```
In [112]: y_pred=model.predict(x_test)
          print(y_pred)
```

```
[ 4.66662368e+03  1.33010678e+04  1.34653152e+04  1.28193039e+04
 1.04531758e+03  3.12026024e+04  1.31016999e+04  1.21308741e+04
 3.76778485e+03  3.02539216e+04  1.18854074e+04  1.72896696e+04
 8.86756871e+03  8.61286966e+03  4.06122202e+03  1.05509036e+04
 4.29383875e+03  6.38320741e+03  1.53904365e+04  1.51357374e+04
 1.25646049e+04  3.24170861e+04  9.24596826e+03  9.94182153e+03
 2.71199666e+03  8.16435467e+03  8.37108330e+03  1.15236009e+04
 7.49432680e+03  4.39165114e+03  1.43900422e+04  5.90331509e+03
 3.32255987e+04  2.72951572e+04  3.29597957e+04  1.00008958e+04
 3.12247474e+04  2.58702645e+04  1.59459962e+04  3.36224001e+04
 6.22444920e+03  1.44712615e+04  1.03183495e+04  1.55787005e+04
 4.05573278e+03  1.31718781e+04  4.86411999e+03  2.95746614e+04
 7.53499914e+03  1.24538798e+04  1.43734491e+04  1.22157738e+04
 2.26535321e+03  8.35817060e+03  2.56229262e+04  1.07797145e+04
 3.37608222e+04  1.52114048e+04  2.16198890e+03  6.75967273e+03
 7.26913354e+03  1.49696184e+04  2.72693318e+04  3.38570492e+03
 1.58075741e+04  1.12098276e+04  1.03404946e+04  1.09477678e+04
 1.24649429e+03  2.50193961e+04  3.66992504e+04  3.23635637e+04
 2.59571964e+03  1.03552788e+04  1.41316627e+04  3.45138782e+04
 3.16606358e+03  4.88445616e+03  1.10954221e+04  9.81631220e+03
-6.80435782e+02  1.34081125e+04  1.01577825e+04  4.02248391e+03
 3.30428239e+04  3.31239805e+04  7.45190830e+03  3.75667119e+04
 1.17838519e+04  9.63908936e+03  3.02354569e+04  3.25611227e+04
 1.45838581e+04  1.11784503e+04  6.18947609e+02  1.16325172e+04
 9.95105385e+03  1.49234568e+04  1.51301855e+04  5.06535939e+03
 1.40855011e+04  2.64092311e+04  2.78046180e+04  2.82050999e+04
 3.62138688e+04  2.72065771e+04  1.27787162e+03  9.77015062e+03
 4.97116470e+03  1.22471511e+04  6.12295643e+03  4.73124988e+03
 1.43475829e+03  1.79375524e+04  3.23805056e+03  2.53296497e+03
 1.04751735e+04  1.28544243e+04  1.01484875e+04  3.79729064e+03
 9.62804819e+03  1.22452796e+04  8.10528039e+03  7.55159223e+03
 3.68930663e+04  1.18373743e+04  1.08665485e+04  2.91058729e+04
 3.59480032e+04  1.16749357e+04  2.90116782e+04  8.55329335e+01
 7.19533771e+03  3.21623870e+04  9.21459093e+03 -8.59906805e+00
 1.56026762e+03  5.28493804e+03  7.85800479e+03  1.23117147e+04
 1.46650774e+04  8.45411145e+03  2.91280179e+04  1.65348047e+04
 1.36554508e+04  1.15568498e+04  2.49235533e+03  9.48220269e+03
 4.31224069e+03  5.55448403e+03  1.19887090e+04  5.04502322e+03
 1.38824529e+04  1.30629618e+04  1.36886370e+04  7.95394564e+03
 1.19905805e+04  1.05472232e+04  9.67233823e+03  5.36247695e+03
 6.67851612e+03  4.06066956e+04  1.37182054e+04  4.76262722e+03
 7.32820782e+03  5.62640832e+03  3.27548760e+04  1.17986362e+04
 1.16915915e+04  6.64345840e+03  6.26137847e+03  6.78924121e+03
 3.31905409e+04  3.51322552e+04  2.35580483e+03  7.03845104e+03
 5.61349562e+03  1.41039658e+04  1.29820779e+03  1.12780715e+04
 1.34856514e+04  1.13612250e+04  1.07170226e+04  1.24722817e+04
 2.14907619e+03  2.85299771e+04  3.12545393e+03  1.50286300e+04
 6.79853622e+03  8.54088267e+03  1.45930904e+04  3.89897958e+04
 3.05165810e+03  1.21511695e+03  4.59101900e+03  8.31943249e+03
 7.05136374e+03  4.57249167e+03  9.51177117e+03  9.37147760e+03
 9.98424003e+03]
```

```
In [113]: diff=pd.DataFrame({"actual":y_test,"predicated":y_pred})  
diff
```

Out[113]:

	actual	predicated
559	1646.43	4666.623676
1087	11353.23	13301.067793
1020	8798.59	13465.315239
460	10381.48	12819.303931
802	2103.08	1045.317580
...
891	7243.81	7051.363739
414	2134.90	4572.491675
258	11520.10	9511.771173
538	8233.10	9371.477595
929	6289.75	9984.240030

201 rows × 2 columns

```
In [114]: from sklearn.metrics import mean_absolute_error  
mae=mean_absolute_error(y_test,y_pred)  
print(mae)
```

4019.341739524793

```
In [115]: from sklearn.metrics import mean_squared_error  
mse=mean_squared_error(y_test,y_pred)  
print(mse)
```

36946790.51390431

```
In [116]: from sklearn.metrics import r2_score  
rs=r2_score(y_test,y_pred)  
print(rs)
```

0.739443551121639

In []: